


CUDA-Accelerated 2D Systolic Array for Low-Bit Quantized Matrix Multiplication

Ben Zhang

Project Motivation:

- Originated from an ASIC project for CNN acceleration using a 2D systolic array.
- Aim: Port and extend the hardware accelerator concept to CUDA for deep learning inference.
- Low-bit (2-bit) quantized neural networks reduce computational and memory requirements.
- Importance of efficient matrix multiplication in CNN inference.



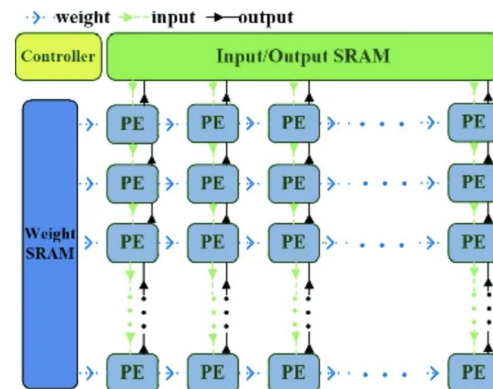
Algorithm

2D Systolic Array Simulation:

- A parallel processing architecture where processing elements (PEs) are arranged in a two-dimensional grid.
- PEs exchange data only with adjacent neighbors.
- Operations are overlapped to maximize throughput.
- Uses local shared storage to reduce global memory accesses.

Low-Bit Quantization:

- Input matrices stored as 8-bit integers but only lower 2 bits are used.
- Optimized arithmetic for quantized data to simulate low-bit operations.



https://link.springer.com/chapter/10.1007/978-3-030-05677-3_16

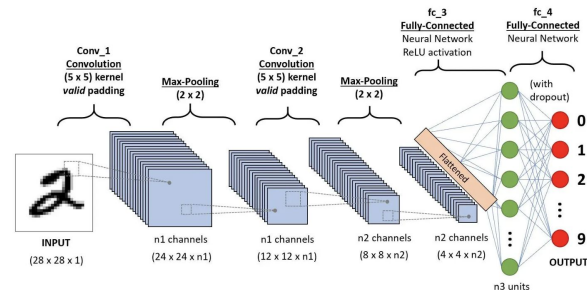
Limitations of Algorithm using CPU

- **Limited Parallelism**
 - CPUs offer significantly fewer cores compared to GPUs, reducing the level of parallel processing available for large-scale matrix multiplications.
- **Performance Bottlenecks**
 - Processing is done sequentially and only in one core, leading to lower throughput.
- **Real-Time Inference Constraints**
 - Achieving low latency and high throughput for deep learning applications is more challenging without the massive parallelism
- **Energy Efficiency**
 - More time to process usually leads to high power consumption

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \end{bmatrix}$$

$$B = \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \\ b_{31} & b_{32} \end{bmatrix}$$

$$C = AB = \begin{bmatrix} a_{11} * b_{11} + a_{12} * b_{21} + a_{13} * b_{31} & a_{11} * b_{12} + a_{12} * b_{22} + a_{13} * b_{32} \\ a_{21} * b_{11} + a_{22} * b_{21} + a_{23} * b_{31} & a_{21} * b_{12} + a_{22} * b_{22} + a_{23} * b_{32} \end{bmatrix}$$



<https://saturncloud.io/blog/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way/>

-



Demo

```
10 feuda_runtime.h
11
12 Microsoft Visual Studio Debug Console
13
14 PASS
15 D:\
16 CPU time: 23.366000 seconds
17 T:\
18 GPU time: 376.973297 ms
19 f\
20 M\
```

Performance and Results

- **Comparative Analysis:**

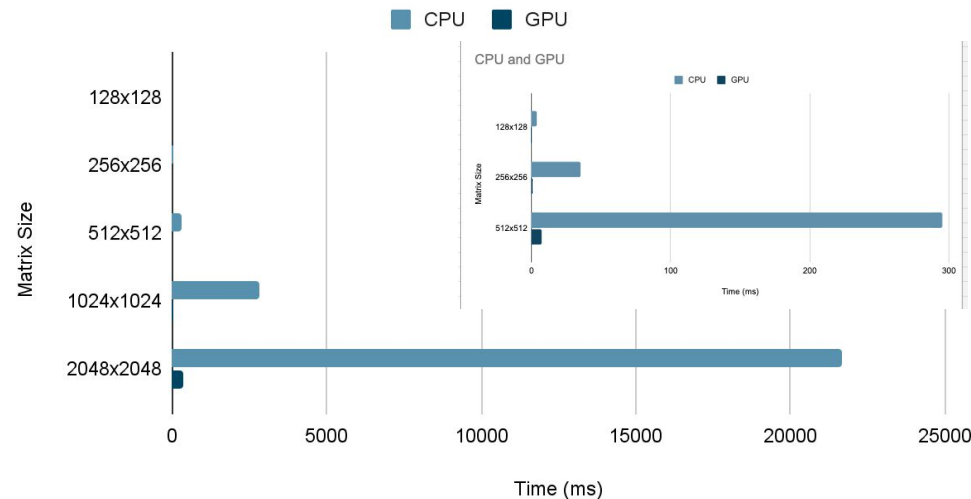
- GPU vs. CPU implementation run times.
- Any Matrix smaller than 128 x 128 will run instantly and CPU timer doesn't work.
- Anything above 2048 x 2048 will take long time for CPU to execute
- Validating correctness by comparing output arrays between CPU and GPU.
- Input matrix are randomly generated with selected Dimension
- $\text{Speedup} = 1/(1-(\text{CPU}/\text{GPU}))$ as assuming "infinite" number of parallel processors for GPU

- **Observations:**

- Significant speed-up using CUDA, especially for larger-scale matrix multiplications.
- Larger Matrix size will have better improvement in execution speed.

Matrix Size	CPU	GPU	Speedup
128x128	4	0.207872	94.80%
256x256	35	0.990208	97.17%
512x512	295	7.084032	97.60%
1024x1024	2836	54.031265	98.09%
2048x2048	22170	353.237	98.41%

CPU and GPU



Conclusion

- **Achievements:**

- Successful design of a CUDA-accelerated systolic array for quantized matrix multiplication.
- Demonstrated performance gains over CPU-based implementations.

- **Future Enhancements:**

- Further memory and kernel optimizations , like use of concurrent kernel, multistream, and even async memory transfer .
- Bit-packing multiple 2-bit values per byte for reduced memory footprint.

References

<https://telesens.co/2018/07/30/systolic-architectures/>

<https://saturncloud.io/blog/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way/>

https://link.springer.com/chapter/10.1007/978-3-030-05677-3_16

https://docs.google.com/presentation/d/1m3y5f82RGP89rW5o2U_iYEZnItNKitjVgIkgyb7-nYY/edit?usp=sharing

<https://github.com/khumpal/ECE284-Final-Project>

Lecture Slides